Mini-Project Proposal

Harnessing AI and Machine Learning

to Supercharge Biotechnology Research

A Hands-On Dive into Protein Structure Prediction and Genomic Analysis

Prepared by: Kajal Singh

Date: June 2026

Exploring the Future of Biotech Innovation

1. Introduction

I've always been fascinated by how biology and technology are colliding in the most exciting ways. From a young age, watching documentaries on genetic engineering and reading about revolutionary medical breakthroughs made me dream of contributing to this field. Biotechnology has given us breakthroughs like CRISPR gene editing and mRNA vaccines that saved millions during the pandemic, but the pace of discovery often feels painfully slow because of the sheer complexity of living systems and the limitations of traditional lab work.

That's where artificial intelligence and machine learning step in as genuine game-changers. These technologies allow us to analyze vast amounts of biological data, uncover hidden patterns, and make predictions that would be impossible for humans alone.

Think about AlphaFold from DeepMind – it basically cracked the 50-year-old protein folding problem, predicting 3D structures from amino acid sequences with stunning accuracy. Suddenly, what used to take months or even years in a wet lab can happen in hours on a computer. This isn't just cool tech; it's opening doors to faster drug development, better understanding of diseases like cancer and Alzheimer's, and even designing brand-new proteins for everything from medicine to sustainable materials and environmental solutions.

In this mini-project, I want to explore AI/ML applications in biotechnology research, focusing on protein structure prediction and genomic analysis for drug discovery. It's a practical, doable project that any enthusiastic student or small team could tackle with open-source tools and publicly available datasets. I'm genuinely excited about how this could help make biotech research more accessible, cost-effective, and impactful for real-world problems.

This proposal reflects my personal passion for interdisciplinary science and my belief that AI can democratize biotechnology, allowing researchers in resource-limited settings to make meaningful contributions.

2. Problem Statement

Let's be honest – traditional drug discovery is brutal. It takes 10-15 years on average and costs billions of dollars, with a heartbreakingly low success rate of around 10% or less for new drugs reaching the market. Scientists spend endless hours in labs screening millions of compounds, determining protein shapes through expensive experimental methods like X-ray crystallography or cryo-EM, and manually sifting through mountains of genomic data.

Meanwhile, diseases like cancer, Alzheimer's, antibiotic-resistant infections, and emerging viral threats keep evolving faster than we can respond. The COVID-19 pandemic highlighted how critical rapid innovation is, yet our processes are still too slow.

Genomic sequencing technologies now generate terabytes of data daily, but turning that raw information into actionable insights – like identifying disease-causing mutations, understanding gene-environment interactions, or predicting how a potential drug will bind to its target protein – remains incredibly challenging for humans alone. The data is messy, massive, heterogeneous, and multidimensional. Traditional methods simply can't keep up with the scale and complexity.

This project zooms in on using AI/ML to tackle these pain points head-on: predicting how proteins fold and interact with potential drugs, and analyzing genomic sequences for useful biomarkers. The ultimate hope is to demonstrate practically how these tools can dramatically speed up discovery, reduce costs by 30-50%, and make biotech innovation more accessible and effective for researchers everywhere.

3. Objectives

Main Goal: To get my hands dirty with AI/ML techniques in biotech and build something that demonstrates real value in accelerating research, particularly around drug discovery. This project is as much about learning and experimentation as it is about producing tangible results.

Specific Things I Want to Achieve:

Dive deep into and test out current AI models (including simplified versions of AlphaFold-like approaches) for predicting protein structures from sequences.

Create a practical ML model that can predict drug-target binding affinities using public datasets like ChEMBL and PDB.

Apply machine learning to genomic data for identifying potential disease biomarkers, using a case study like cancer-related genes.

Evaluate model performance using metrics such as accuracy, precision, and comparison with traditional lab benchmarks.

Think critically about real-world scalability, limitations, and the ethical dimensions including data privacy, model bias, and ensuring AI augments rather than replaces human expertise in biotechnology.

4. Methodology

I plan to keep this grounded and iterative – more like a real research journey than a rigid plan. The project will be hands-on, using freely available tools to ensure reproducibility and accessibility.

Phase 1: Getting Grounded (Weeks 1-2) - Conduct a thorough literature review of key papers on AlphaFold, RoseTTAFold, and ML in genomics. Collect and preprocess open datasets from Protein Data Bank (PDB), UniProt, ChEMBL, and NCBI genomic resources.

Phase 2: Building the Models (Weeks 3-6) - Implement and fine-tune models for protein structure prediction using graph neural networks or pre-trained models. Develop supervised ML for drug-target interactions and apply sequence analysis techniques (CNNs or transformers) for genomic variant detection. Tools include Python, PyTorch, scikit-learn, BioPython, and RDKit.

Phase 3: Testing and Integration (Weeks 7-8) - Train models with proper validation, tune hyperparameters, and create a user-friendly interface using Streamlit for easy interaction.

Phase 4: Evaluation and Reflection (Weeks 9-12) - Assess performance with standard metrics, compare against baselines, discuss limitations honestly, and explore ethical implications like data privacy and bias mitigation.

5. Expected Outcomes

By the end of this project, I expect to have a functional AI/ML prototype capable of predicting protein structures for several examples, screening compound libraries for potential drug candidates, and identifying genomic biomarkers.
Deliverables will include clean, well-documented code on GitHub, a comprehensive report with visualizations and results, and a presentation that I can share with peers or mentors.
More importantly, this work will provide personal insights into how AI can accelerate biotech research, potentially highlighting promising leads for further study. On a broader scale, it contributes to the vision of making advanced tools accessible, fostering innovation in healthcare, agriculture, and climate solutions.
Success will be measured not just by model accuracy but by the learning experience and potential for future extensions, such as collaboration with wet-lab teams.

6. Timeline, Budget & References

Timeline: 10-12 weeks for a complete mini-project cycle. Budget: Very minimal – under $300 if cloud GPU resources are needed; otherwise entirely free using open-source tools and public datasets.
Key References:
1. Jumper, J., et al. (2021). Highly accurate protein structure prediction with AlphaFold. Nature.
2. Recent reviews (2024-2026) on AI/ML applications in drug discovery and genomics from journals like Nature Biotechnology and PMC.
Additional sources will be cited properly in the final report.
This proposal is written with genuine enthusiasm and personal reflection on the transformative potential of AI in biotechnology. I look forward to embarking on this journey!